

Rhazel Anne Galura

Senior Fullstack AI Engineer

Philippines (GMT+8) | Remote / Distributed Teams

galurarhazel7@gmail.com | <https://www.linkedin.com/in/rhazel-g-7127013b9/>

Summary

"Build things that actually work in the real world."

That's how I approach AI.

Hands-on AI fullstack engineer who builds AI products that actually ship. Six-plus years turning messy data, product ideas, and business needs into production-ready AI systems for US/EU enterprises and early-stage startups, across backend services, APIs, data pipelines, LLM and agent workflows, model deployment, and production reliability, with a focus on regulated domains like healthcare and finance.

Experience

Senior AI Fullstack Engineer / Azure Fabric Engineer — Remote

QTX / Accenture

Oct 2025 – Jun 2026

GenAI Assistant & Data Platform (Aged Care)

- Built a production-grade GenAI assistant PoC for BaptistCare (Australian aged care): a RAG system over clinical and policy documents on Azure OpenAI and Azure AI Search, with a FastAPI backend and a streaming React/Next.js UI.
- Designed the document-intelligence pipeline (Azure Document Intelligence, semantic chunking, embeddings, hybrid retrieval) and an evaluation framework measuring retrieval recall, answer faithfulness, and latency.
- Led the Whiddon data platform: an Azure Synapse medallion warehouse, Azure Data Factory ingestion across 5-7 source systems, master data management, and Power BI executive reporting with row-level security.
- Built predictive analytics for resident deterioration risk and funding/staffing forecasts, and set the architecture for event-driven real-time integration via Azure Service Bus.

Senior AI Fullstack Engineer — Remote

Accenture

Jan 2020 – Jun 2026

Enterprise, Healthcare Startup (0 → 1 → Scale) & AI / Data Consulting

- Led production ML/AI systems for a Fortune 500 US financial-services client, improving fraud detection accuracy by +18%.
- Built and deployed LLM applications (chat assistants, document QA) using fine-tuned models and RAG pipelines, plus agent-based workflows for enterprise automation.
- Fine-tuned LLMs on domain data and built prompt-engineering frameworks to improve accuracy and reduce hallucinations and cost.
- As primary ML/AI lead for a healthcare startup, built core AI features 0→1 (clinical/document assistants, recommendation systems, CV medical-document pipelines), cutting model latency ~30% and scaling to tens of thousands of users.
- Delivered AI across fintech, e-commerce, and logistics: NLP systems, RAG and fine-tuning pipelines, and computer-vision automation, from pilot to production.

Education

Master of Engineering (MEng) in Artificial Intelligence

University of the Philippines Diliman

2019 – 2021

Bachelor of Science (BS) in Computer Science

University of the Philippines Diliman

2016 – 2019

Additional

- 6+ years of experience building and deploying end-to-end AI systems in production
- Extensive experience working with distributed teams across US, EU and APAC
- Strong focus on shipping reliable, production-grade AI products across backend services, APIs, data pipelines, and ML systems
- Regulated-domain delivery in healthcare, aged care and financial services

Skills

AI / Machine Learning

- PyTorch
- TensorFlow
- scikit-learn
- NLP
- computer vision (OCR, YOLO, OpenCV)
- recommendation systems

Agentic & Advanced AI Systems

- LLM fine-tuning (LoRA)
- RAG pipelines
- embeddings
- prompt engineering
- agent architectures
- LangChain-style orchestration
- vector databases

Frameworks & Libraries

- Hugging Face Transformers
- Azure OpenAI

Cloud, MLOps & Infrastructure

- Azure (Synapse, Data Factory, AI Search)
- AWS (S3, EC2, Lambda, SageMaker)
- GCP (Vertex AI, BigQuery)
- Docker
- CI/CD (GitHub Actions)
- model deployment & monitoring

Data Engineering & Systems

- SQL
- ETL/ELT
- Power BI

Software Engineering

- Python
- Java
- FastAPI
- REST APIs
- microservices
- React