

Rhazel Anne Galura

Senior Fullstack AI Engineer

Philippines (GMT+8) | Remote / Distributed Teams | Open to Global Roles
galurarhazel7@gmail.com | <https://www.linkedin.com/in/rhazel-g-7127013b9/>

Summary

“Build things that actually work in the real world.”

That’s how I approach AI.

I’m a hands-on AI Fullstack Engineer focused on building AI products that actually ship. Over the past 6+ years I’ve worked with US, EU and APAC enterprise teams and early-stage startups to turn messy data, product ideas, and business needs into production-ready AI systems. My work spans the full stack of applied AI: backend services, APIs, data pipelines, LLM workflows, agent systems, model deployment, and production reliability.

I work at the intersection of engineering, product, and AI — designing LLM-powered features, building agent workflows, developing application backends, integrating external tools and data systems, and deploying reliable AI platforms in production. I’m especially interested in LLMs and applied AI in regulated domains like healthcare and finance.

Experience

Senior AI Fullstack Engineer / Azure Fabric Engineer — Remote

QTX / Accenture

Oct 2025 – Jun 2026

GenAI Assistant (Aged Care)

- Built the production-grade GenAI assistant for BaptistCare, an Australian aged care provider, as a retrieval-augmented generation system over clinical and policy documents, taking it from proof of concept to a working product used by staff
- Implemented the assistant on Azure OpenAI and Azure AI Search, with a FastAPI backend and a streaming React / Next.js interface that returned answers progressively with citations back to the source document
- Designed the document-intelligence pipeline using Azure Document Intelligence, semantic chunking and embedding generation, then layered hybrid retrieval over it so keyword and vector matches were combined before ranking
- Built the evaluation framework that measured retrieval recall, answer faithfulness and latency, and used it to choose chunking strategy, retrieval configuration and prompt structure on evidence rather than intuition
- Added guardrails around clinical content so the assistant grounded every answer in retrieved policy text and declined rather than guessing when confidence was low, which is the behaviour a regulated care setting requires
- Worked directly with clinical and operational stakeholders to turn their questions into evaluation cases, then iterated retrieval and prompting against those cases through successive releases

Data Platform Modernization & Predictive Analytics

- Led the Whiddon data platform modernization, architecting an Azure Synapse medallion warehouse with bronze, silver and gold layers and defining the contracts between them
- Built Azure Data Factory ingestion across 5-7 source systems, with master data management to reconcile resident, staff and facility records that had diverged between operational systems
- Delivered Power BI executive reporting on top of the warehouse with row-level security, so leadership, facility managers and clinical teams each saw only the data their role permitted
- Built predictive analytics for resident deterioration risk and for funding and staffing forecasts, operationalizing the models into production scoring jobs that fed care and rostering workflows rather than sitting in notebooks
- Set the architecture for event-driven real-time integration using Azure Service Bus, streaming updates between clinical, operational and analytics systems with schema contracts, poison-queue handling and retry with backoff
- Owned feature lifecycles end to end across both workstreams, authoring API contracts, deployment pipelines and the production runbooks used during releases
- Instrumented the platform for monitoring and incident response, and took part in on-call for the assistant and data pipelines
- Mentored engineers on component testing and CI/CD practice, and introduced trunk-based development with GitHub Actions to shorten the path from merge to release

Senior AI Fullstack Engineer / AI Consultant — Remote

Accenture (Contract, Remote)

Jan 2020 – Oct 2025

Enterprise

- Led development of production-grade ML and AI systems for a Fortune 500 US financial services client, improving fraud detection accuracy by +18%
- Built and deployed LLM-powered applications including chat assistants and document QA systems using fine-tuned models and RAG pipelines
- Designed and implemented backend AI services, API orchestration, and agent-based workflows for enterprise automation
- Fine-tuned large language models on domain-specific datasets to improve response accuracy and reduce hallucinations
- Developed computer vision pipelines for document processing and visual inspection
- Implemented prompt engineering frameworks to improve reliability, consistency, and cost efficiency of LLM systems
- Built scalable ML pipelines spanning training, deployment, monitoring, and production operations
- Collaborated with US-based stakeholders across product, engineering, and business teams to deliver production-ready AI solutions

Startup (0 -> 1 -> Scale) — Consulting Engagement, Healthcare

- Led AI development for an early-stage healthcare technology startup client, serving as the primary ML/AI lead from 0 to first production release
- Designed and built core AI product features from scratch, including:
 - LLM-powered clinical and document assistants for summarization, extraction, and knowledge workflows
 - Recommendation systems for workflow prioritization and decision support
- Developed application backends, APIs, agent workflows, data pipelines, and frontend-integrated AI features to automate healthcare operations
- Worked across the product stack to connect AI systems into user-facing workflows and internal tools, supporting fast iteration from prototype to production
- Contributed light frontend work in React to help integrate AI capabilities into product features and internal applications
- Built workflows for patient intake, document processing, and clinical / administrative decision support
- Delivered a platform supporting real-time inference and early growth from 0 to tens of thousands of users / transactions
- Built computer vision pipelines for medical document processing, including forms, records, and structured data extraction
- Reduced model latency by ~30%, enabling reliable real-time AI features
- Established foundational data infrastructure across ETL, APIs, cloud systems, and product integrations with secure data handling aligned with healthcare requirements
- Partnered with product and engineering teams to ensure AI systems met usability, accuracy, and compliance needs
- Supported the transition from prototype to production with monitoring, evaluation, and reliability mechanisms

AI / Data Consulting Work

- Delivered AI solutions across fintech, e-commerce, and logistics
- Built:
 - NLP systems including chatbots, classification pipelines, and embeddings-based applications
 - LLM pipelines including RAG, fine-tuning, and prompt chaining
 - Computer vision systems for automation and document processing
- Designed full-stack AI architectures integrating LLMs with APIs, databases, external tools, and production services
- Advised clients on AI system design, architecture, deployment, scaling, and evaluation best practices
- Supported end-to-end delivery from pilot to production, ensuring performance, reliability, and maintainability

Education

Master of Engineering (MEng) in Artificial Intelligence

University of the Philippines Diliman

2019 – 2021

Bachelor of Science (BS) in Computer Science

Additional

- 6+ years of experience building and deploying end-to-end AI systems in production
- Extensive experience working with distributed teams across US, EU and APAC
- Strong focus on shipping reliable, production-grade AI products across backend services, APIs, data pipelines, and ML systems
- Regulated-domain delivery in healthcare, aged care and financial services

Skills

AI / Machine Learning

- Machine Learning, Deep Learning, Applied AI
- Generative AI (LLMs, multimodal models)
- LLM fine-tuning (LoRA, instruction tuning), RAG pipelines, embeddings
- Prompt engineering, prompt chaining, evaluation & guardrails
- NLP (classification, NER, semantic search, text similarity)
- Computer Vision (image classification, object detection, OCR, document AI)
- Recommendation systems, ranking models
- Time-series forecasting, predictive risk scoring

Agentic & Advanced AI Systems

- Agent-based architectures (multi-step reasoning, tool use, memory)
- LLM orchestration frameworks (LangChain-style pipelines, custom agents)
- Retrieval systems (vector DBs, hybrid search)
- Decision-support systems & workflow automation
- Real-time AI inference systems

Frameworks & Libraries

- PyTorch, TensorFlow, Scikit-learn
- Hugging Face Transformers, Diffusers
- OpenAI API / LLM APIs, Azure OpenAI
- OpenCV, YOLO, Tesseract (OCR), Azure Document Intelligence

Cloud, MLOps & Infrastructure

- AWS (S3, EC2, Lambda, SageMaker), GCP (Vertex AI, BigQuery)
- Azure (Synapse, Data Factory, AI Search, Service Bus)
- Docker, containerized deployments
- CI/CD pipelines (GitHub Actions, GitLab CI)
- Model deployment, monitoring, observability, experiment tracking
- ML lifecycle: training -> validation -> deployment -> monitoring -> iteration
- Batch + real-time processing systems

Data Engineering & Systems

- SQL, data modeling, data warehousing, medallion architecture
- ETL/ELT pipelines, data ingestion, master data management
- Feature engineering, feature stores
- APIs, event-driven pipelines
- Data integration across cloud and external systems
- Power BI reporting with row-level security

Software Engineering

- Python (primary), Java
- FastAPI, REST APIs, microservices

- React, Next.js
- System design for scalable AI systems
- Distributed systems basics
- Performance optimization & latency reduction

Product & Collaboration

- Translating business requirements -> AI solutions
- Working with cross-functional teams (product, design, engineering)
- Stakeholder communication (especially US clients)
- Experimentation, A/B testing mindset
- Agile / iterative development